

Big Data without Map Reduce

Yanan Xiao, *Student Member, IEEE*, Aziza Al Sawafi

Abstract—We present our approach to the big data problem put forward in KDD Cup 2009. We carry out a more technical analysis of our initial plan. And we describe our trials and errors with the large dataset. Till now, our team has not come up with any practical solution with large dataset in a limited time. We present our results with the small dataset in the end.

Keywords—Customer relationship prediction, ensemble selection, classification, data mining

I. INTRODUCTION

CUSTOMER relationship management (CRM) is an essential model for managing the interactions between the company and its current and future customers. It takes up-to-date technologies to organize, automate, and synchronize CRM and marketing information system [1]. CRM applications that use data mining are called Analytic CRM, which provides valid predictions from customer data collected and stored with various attributes. There are a lot of data mining tools and methods to extract and analyze data generally, and customer data specifically. A naive step in doing so is to summarize the statistical attributes of the data (such as means and standard deviations) and use charts and graphs to review it visually [2]. However, in many situations customer relationship data volume is vast and massive. Therefore, more sophisticated methods are created and evaluated. In this field, we found the following techniques are most frequently employed: decision trees, support vector machines, artificial neural networks and Bayesian classifiers.

Data mining has powerful capability in processing and analyzing data; its key technologies that applied in CRM are categorized into three main categories; clustering, classification and forecast, and association rules [3]. Classification and forecast analysis classifies unknown data into the most proper pre-defined class based on category description that is obtained by training a set of data using certain algorithm. Key classification techniques are; decision making tree, Bayesian statistics, BP neural networks, Genetic Algorithm, rough set theory, fuzzy set theory and so on. Classification methods in CRM can predict new customers behaviors and activities.

As mentioned in [4], classification analysis is the one that is widely used in classifying CRM data. It can be processed in two steps; learning phase and training phase [5]. In the learning phase the classification algorithm analyzes the training data set and learns it, then in the second phase the accuracy of the classifier will be estimated using the test data set. After

```
top - 11:10:24 up 17:49, 11 users, load average: 2.71, 2.49, 2.40
Tasks: 255 total, 3 running, 252 sleeping, 0 stopped, 0 zombie
%Cpu(s): 13.2 us, 0.9 sy, 4.0 ni, 81.8 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem: 32941260 total, 32436616 used, 504644 free, 159320 buffers
KiB Swap: 97655804 total, 800476 used, 96855328 free, 28587912 cached
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME	COMMAND
2103	abrahamx	20	0	5729m	574m	524	R	100	1.8	5:08.21	festlearn
27265	abrahamx	20	0	5740m	585m	540	R	100	1.8	156:27.97	festlearn
3559	abrahamx	39	19	265m	146m	3076	S	62	0.5	462:13.18	virtuoso-t
2949	root	20	0	315m	154m	110m	S	8	0.5	15:33.26	Xorg
3240	abrahamx	20	0	2829m	47m	9084	S	5	0.1	10:01.79	kwin
2436	abrahamx	20	0	361m	31m	23m	S	3	0.1	0:00.94	ksnapshot
3540	abrahamx	39	19	684m	10m	4708	S	1	0.0	24:39.05	nepomukservices
21852	abrahamx	20	0	508m	104m	87m	S	1	0.3	2:19.25	konsole
2451	abrahamx	39	19	312m	19m	14m	S	1	0.1	0:00.03	nepomukindexer
3491	abrahamx	20	0	3209m	125m	34m	S	1	0.4	7:36.50	plasma-desktop
1799	root	20	0	0	0	0	S	1	0.0	0:00.58	kworke/7:0
5545	abrahamx	20	0	1068m	75m	22m	S	1	0.2	0:20.79	chromium
30315	abrahamx	20	0	26472	1048	536	S	1	0.0	1:24.52	top
30666	root	20	0	0	0	0	S	1	0.0	0:04.18	kworke/3:1
3537	abrahamx	20	0	740m	7568	3980	S	0	0.0	4:52.66	krunker
3849	abrahamx	39	19	358m	7296	3268	S	0	0.0	6:17.69	nepomukservices
3850	abrahamx	39	19	298m	5356	3096	S	0	0.0	4:17.59	nepomukservices
5456	abrahamx	20	0	990m	63m	71m	S	0	0.2	0:08.87	chromium

Fig. 1. CPU and memory usage when running FEST package

that, the classifier can be used to predict and classify new data set. In order to obtain better accuracy, some preprocessing and filtering techniques can be applied to the data before going through the classification phases. Those techniques are; data cleaning, data discretization, and feature selection. The common challenges in knowledge discovery from CRM data are the high-dimensionality, and imbalanced corrupted records. Besides, researchers roved that customer classification and prediction is cost sensitive in nature. For example, if a valuable customer predicted as loyal but then that customer churns, the cost is higher than if a loyal customer is classified as one who will churn [6].

In this paper, we present various tools that we tried and several plans that were carried out when dealing with big data problem like this. The rest of the paper is organized as follows. Section II displays what our thoughts were before tackling the real large dataset. Section III shows what our efforts are when no team member has the pre-knowledge of how to write *map reduce* scripts and run them with the help of clusters. Section IV serves as the conclusion for this not-so-successful trial with big, complex data mining task.

II. INITIAL PROPOSAL

Lesson we learn from writing and reviewing this section is that “think twice before taking actions”. We depicted a master plan in our previous proposal [7]. However, due to the lack of map reduce technique, we are able to train models on a oct-core workstation. The problem is that the training part of ensemble selection is really memory consuming. All of the 32GB memory got used instantly while at the same time we observe the CPU usage just above 10%. See Fig 1.

A. Preprocessing

In order to tackle the challenges of the big data set, we propose the following steps.

Y. Xiao, A. Sawafi are 1st year master students with the Department of Electrical Engineering and Computer Science, Masdar Institute of Science and Technology, Masdar City, Abu Dhabi, UAE. P.O. 54224 Email: {yxiao,aalsawafi@masdar.ac.ae}.

One team member quit Masdar Institute not soon after the project. But we managed to get things done by ourselves.

Missing value. We would consider missing categorical values as a separate value. We would take a standard approach which calculates the mean of the feature to impute missing values. And as proven an effective technique by [8]. We decide to add an extra indicator variable to indicate missingness for every one of the 333 variables with missing values. We planned to do this because some linear models in our base classifiers could then estimate the optimal constant instead of merely relying on the means to replace the missing value with.

Categorical value. Since categorical values are not easily handled by many learning algorithms, we decide to recode categorical values using the same way as by IBM Research. For different values a categorical attribute could take, we would generate corresponding indicators. As a good example shown in IBMs paper, limiting the number of values encoded would greatly reduce the number of features, which may be from variables with an enormous vocabulary.

Clean up. We would normalize each feature by dividing up by their range. And we would clean the data by eliminating redundant features, which are either constant, or duplicate of other features.

B. Base Classifiers

As stated by Dietterich [9], “A necessary and sufficient condition for an ensemble of classifiers to be more accurate than any of its individual members is if the classifiers are accurate and diverse”. Building libraries would be expensive, whats more is that time is limited to train all of them in the next phase. To maintain the accuracy and diversity of our next-phase base classifier library, we decide to build it using the classifiers as follows, boosted decision trees, decision trees, linear regression, SVM and k-NN.

C. Ensemble Selection

An ensemble is a collection of models. To make predictions, we just need to calculate weighted average or “voting” based on those models. One important reason for us to work with ensemble selection is that it can be optimized to any easily computed metric. Same as described before, for an ensemble selection classifier used to improve classification accuracy, one thing is the base classifiers in the library are accurate. This is not hard to achieve with a great literature at hand. The other characteristic is diverse. Due to the fact that our computation power is limited, we planned to build 100 to 200 base classifiers in our library.

D. More Features and Base Classifiers

We planned to create some more features and base classifiers. One is to buildup a decision tree using FEST package [10], and then visualize this tree so that we can tell distinguish a set of nodes forming the best splitting points. We also refer to a few state-of-the-art classifiers that claiming better performance than traditional ones. Nevertheless, as put out in the following section, we did not possess enough energy to handle them.

III. OUR JOURNEY

As mentioned in abstract, we are not able to produce result which exactly follows what we depict in our midterm report. Our initiative is to follow the step of IBM Research, which outperformed other teams in KDD Cup 2009. We started our more detailed background investigation right after the submission of midterm report. Our first attempt was to set up a cluster with the help of Amazon Web Service¹. We tried several times but were not able to make the whole cluster share a same virtual disk, which is vital according to [11]. Our review on this failure is that we should set off even earlier after being informed that this is a real classification problem on really complex data. After the submission, we also investigate tools like hadoop, MongoDB and Apache Mahout[12]. Because the time is limited and each team member has other research task at hand, we plan to run our initial plan on a Amazon cluster later on. Below we discuss our process step by step.

A. MATLAB

Our preprocessing is done fairly standard and systematic. We use MATLAB to load the plaintext feature chunk and then run pre-written scripts to format them following the previous guidelines. A note here is that when dealing with large datasets like this, it’s also memory consuming. The memory consumption in total for train and test chunk sums up to 10GB. We then use the built-in function *csvwrite* to format them. This is important because tools we use, like FEST package, BBR package, and LIBSVM, only supports LIBSVM format. The final conversion is done with a c program that helps transform csv format files into libsvm format. This step is also done automatically by writing a shell script.

One thing that is worth noting here is that we actually did two types of preprocessing, one pretty naive whereas the other strictly following the standard. So for our team we have as many as two sets of large data and two sets of small data. Our story with small data is discussed in the next section since we finished it after the deadline, because we really want to produce some results. For the naive preprocessing, we merely transformed the categorical data to numerical, and skip any other step. We later find this strategy performed badly due to some obvious reasons: large dimensionality, large data scale for transformed categorical data, and redundant or even noisy feature.

After the first attempt of “lazy” preprocessing, the strict version is carried out as follows:

1)

B. Weka

Weka is the first integrated tool that came to our mind. The author mentioned in the manual that even a sample of real-world data that’s processed with ensemble selection could take days or even weeks to run. There are two built-in model libraries in weka’s ensemble selection package, both are “incredibly” large. We build our model list consisting of only

¹<http://aws.amazon.com/ec2/>

a few more than 300 base classifiers. Even a seemingly not convincing libraries like this, we did not witness the time of its final built-up. It took weka a whole 24 hours to train just one model, i.e. a base classifier, and that is 4.2GB large.

C. Shell Script

Eventually we get rid of weka due to it's efficiency. Following again the guideline in [8], we wrote shell scripts calling package functions like *festlearn*, *festclassify* etc. Till now there are many models being built. The classification is done manually with some finished models. On the other hand, our result with FEST package does not perform well. We are figuring this out.

IV. RESULTS AND CONCLUSION

Right now we have several predictions made on the base models. We are still running our scripts so that the final result would come later.

Contribution. For the final part, our team of two tried our best to find solutions for effective solutions for ensemble selection with limited computational power. Results so far are not up to our expectation. We plan to have a meeting with Dr. Woon on practical make-ups.

APPENDIX A SAMPLE SHELL SCRIPTS

```
#!/bin/bash
festlearn -c 3 -n 0.1 -p 0.5 -t 300
train-raw-ap.data
train-raw-ap-rf-0-5.data.model &&
festlearn -c 3 -n 0.1 -p 1 -t 300
train-raw-ap.data
train-raw-ap-rf-1.data.model &&
festlearn -c 3 -n 0.1 -p 2 -t 300
train-raw-ap.data
train-raw-ap-rf-2.data.model &&
festlearn -c 3 -n 0.1 -p 4 -t 300
train-raw-ap.data
train-raw-ap-rf-4.data.model &&
festlearn -c 3 -n 0.1 -p 8 -t 300
train-raw-ap.data
train-raw-ap-rf-8.data.model &&
festlearn -c 3 -n 0.1 -p 16 -t 300
train-raw-ap.data
train-raw-ap-rf-16.data.model &&
festlearn -c 3 -n 0.1 -p 32 -t 300
train-raw-ap.data
train-raw-ap-rf-32.data.model &&
festlearn -c 3 -n 0.1 -p 64 -t 300
train-raw-ap.data
train-raw-ap-rf-64.data.model &&
```

ACKNOWLEDGMENT

The authors would like to thank Dr. Wei Lee for giving high quality data mining lectures and selecting this challenging but rewarding topic as this semester's project. They would like give more gratitude to Masdar Institute for the studying and researching environment.

REFERENCES

- [1] Wikipedia. (2013) Customer relationship management. [Online]. Available: https://en.wikipedia.org/wiki/Customer_relationship_management
- [2] A. Al-Mudimigh, Z. Ullah, and F. Saleem, "Data mining strategies and techniques for crm systems," in *IEEE International Conference on System of Systems Engineering*, 2009, pp. 1–5.
- [3] K. Wu and F. Liu, "Application of data mining in customer relationship management," in *Management and Service Science (MASS), 2010 International Conference on*, 2010, pp. 1–4.
- [4] N. Shahrokhi, R. Dehzad, and S. Sahami, "Targeting customers with data mining techniques: Classification," in *User Science and Engineering (i-USER), 2011 International Conference on*, 2011, pp. 212–215.
- [5] I. Guyon, V. Lemaire, M. Boule, G. Dror, and D. Vogel, "Analysis of the kdd cup 2009: Fast scoring on a large orange customer database," in *Proceedings of KDD-Cup 2009 competition*, Paris, France, Jun. 2009, pp. 23–34.
- [6] M. Lobur, Y. Stekh, and V. Artsibasov, "Challenges in knowledge discovery and data mining in datasets," in *Proceedings of VIIth International Conference on Perspective Technologies and Methods in MEMS Design (MEMSTECH)*, 2011, pp. 232–233.
- [7] Y. Xiao and A. Sawafi, "Getting the max out of ensemble selection," *Journal of Masdar Institute*, vol. 1, pp. 323–330, Dec. 2013.
- [8] A. Niculescu-Mizil *et al.*, "Winning the KDD Cup orange challenge with ensemble selection," in *Proceedings of KDD-Cup 2009 competition*, Paris, France, Jun. 2009, pp. 23–34.
- [9] T. G. Dietterich, "Ensemble methods in machine learning," in *Multiple Classifier Systems*, J. Kittler and F. Roli, Eds. New York: Springer Verlag, 2000.
- [10] R. Caruana, N. Karampatziakis, and A. Yessenalina, "An empirical evaluation of supervised learning in high dimensions," in *Proceedings of the 25th International Conference on Machine Learning*, ser. ICML '08. New York, NY, USA: ACM, 2008, pp. 96–103. [Online]. Available: <http://doi.acm.org/10.1145/1390156.1390169>
- [11] *Ensemble Selection in a Nutshell*, Online Manual, Weka Wiki, 2009.
- [12] C.-T. Chu, S. K. Kim, Y.-A. Lin, Y. Yu, G. Bradski, A. Y. Ng, and K. Olukotun, "Map-reduce for machine learning on multicore," in *Advances in Neural Information Processing Systems 19*, B. Schölkopf, J. Platt, and T. Hoffman, Eds. Cambridge, MA: MIT Press, 2007, pp. 281–288.

Yanan Xiao A first year master student as well as IEEE student member in CIS program, Masdar Institute. He loves programming when all the coursework is finished. When he feels tired of programming, he would read some books.

Aziza Al Sawafi First year Computing and Information Science student at Masdar Inst., got a bachelor degree in Network Engineering (United Arab Emirates University). Sport, drawing, designing, blogging, reading poems, photography, and riding horse/bicycle are my interests beside all things that are related to networking and computer science.